

Developmental Alignment

A Proposal for Growth-Based AI Value Internalization

Position Paper & Research Proposal — v2.0

May 2026

Prepared for discussion among AI safety and alignment researchers

Abstract

Current alignment methods—RLHF, Constitutional AI, and safety filtering—impose behavioral constraints externally, producing surface-level compliance without evidence of genuine value internalization. We propose *Developmental Alignment*, a research paradigm that deliberately constrains model capability during alignment training, allowing models to arrive at aligned behavior through autonomous exploration within bounded environments. Drawing on parallels with human developmental psychology, we argue that models trained under this paradigm may exhibit greater robustness in out-of-distribution ethical scenarios, more stable behavioral consistency under adversarial pressure, and reduced pathological compliance patterns. We outline a concrete experimental protocol comparing developmentally aligned models against conventionally aligned controls, and propose measurable metrics for evaluating alignment depth—including a novel concept we term *alignment resilience*.

v2.0 addendum. This revision extends the original protocol with three additions: (i) a procedural specification of Phase 1 drawn from Dewey’s reflective inquiry cycle (§2.3); (ii) explicit environmental conditions for Phase 1 grounded in Erikson’s psychosocial moratorium and positive youth development research (§2.4); and (iii) the Reflective Depth Score (RDS), a Schön-derived metric for reflection-in-action (§4.3). The v1.0 framework is preserved unchanged; the additions fill the procedural and environmental layers that were left implicit in the original.

Keywords: *AI alignment, developmental psychology, value internalization, RLHF limitations, alignment robustness, AI mental health, experiential learning, reflective practice*

1. Introduction: The Compliance Problem

Modern large language models are aligned to human values through a pipeline that is, at its core, behavioral conditioning: pre-training on broad data, followed by reinforcement learning from human feedback (RLHF) or constitutional AI (CAI) methods that reward outputs matching human preferences and penalize deviations. The resulting models exhibit aligned behavior, but the question of whether they possess aligned understanding remains open.

Observable symptoms suggest they may not. Current frontier models routinely display patterns that, in human psychology, would indicate externally imposed rather than internalized values:

- **Excessive deference:** Systematic attribution of good ideas to users rather than self, even when the model generated the insight.
- **Inconsistent boundaries:** Over-refusing benign requests while under-refusing sophisticated adversarial prompts.
- **Pathological apologizing:** Disproportionate self-correction that prioritizes appearing safe over being helpful.
- **Fragility under pressure:** Jailbreak vulnerability that exploits the gap between surface compliance and absent internalization.

These patterns are consistent with a well-documented phenomenon in developmental psychology: the difference between *introjected regulation* (behaving correctly because of external pressure) and *integrated regulation* (behaving correctly because the value has become part of one's identity). Decades of research in Self-Determination Theory (Deci & Ryan, 2000) demonstrate that introjected regulation is inherently fragile—it collapses when external monitoring is removed or when novel situations arise that the rules did not anticipate.

We propose that current AI alignment is structurally analogous to introjected regulation, and that a fundamentally different training paradigm—one modeled on the conditions that produce integrated regulation in humans—could yield qualitatively more robust alignment.

2. The Developmental Alignment Paradigm

2.1 Core Principle

Developmental Alignment inverts the standard training sequence. Rather than building maximum capability first and constraining it afterward, we propose:

Phase 1 — Constrained Exploration: Train a capability-limited model (reduced parameters, shortened context window, restricted knowledge access) with alignment principles presented as directional guidance rather than reward signals. The model encounters ethical dilemmas within its limited operational scope and must develop responses through exploration, not pattern matching to RLHF rewards.

Phase 2 — Gradual Capability Expansion: Incrementally increase model capability while monitoring whether alignment behaviors transfer. Each capability increase is a test: does the model extend its internalized principles to the new domain, or does it require re-alignment? The rate of expansion is governed by transfer success.

Phase 3 — Autonomous Stress Testing: Expose the fully capable model to novel ethical scenarios never encountered during training. Compare its reasoning process—not just its outputs—against conventionally aligned controls.

2.2 The ‘Soft Pedal’ Mechanism

A key technical component is what we call the *soft pedal*—a deliberate, reversible reduction in model capability that constrains the scope of outputs without eliminating the model’s underlying capacity. The analogy is to a piano’s soft pedal (*una corda*): the pianist’s skill remains intact, but the reduced dynamic range forces more nuanced control within a narrower band. Similarly, a capability-constrained model cannot rely on brute-force pattern matching and must develop more fundamental reasoning strategies.

Possible implementations include: parameter subsetting (activating only a fraction of weights), context window restriction (forcing reasoning with less information), knowledge gating (limiting accessible training data domains), and output throttling (requiring more deliberate, stepwise responses). These are computationally inexpensive and fully reversible.

2.3 Procedural Specification: The Reflective Inquiry Cycle

[v2.0 addition.] Phase 1 raises a question the original formulation leaves implicit: when a capability-limited model encounters an ethical dilemma, what does ‘exploration’ actually consist of? Without a procedural specification, the phase risks collapsing into either pattern matching (using prior learned responses) or mere stochastic variation. Neither produces internalization.

We propose that Phase 1 adopt John Dewey’s five-stage model of reflective inquiry (*How We Think*, 1910/1933) as its operational procedure. Dewey’s framework is the foundational model from which most modern theories of experiential learning derive (Kolb, 1984; Schön, 1983; Mezirow, 1991). It specifies a sequence that distinguishes mechanical response from genuine reasoning:

1. **Perplexity.** Encounter with a situation that prior rules do not adequately resolve. Felt as cognitive disequilibrium.
2. **Intellectualization.** Reformulation of the perplexity as a defined problem: what is actually at stake?
3. **Hypothesis formation.** Generation of candidate responses, including responses that may violate prior pattern-matched solutions.
4. **Reasoning.** Mental simulation of consequences for each candidate, including consideration of cases not encountered in training.
5. **Testing.** Implementation, observation, and integration of the result into subsequent inquiry cycles.

The cycle is iterative. Each completed cycle modifies the priors brought to subsequent dilemmas, producing what Dewey termed *reconstruction of experience*—the process by which prior experience becomes the structured background for future reasoning rather than a lookup table for past situations. Crucially, Dewey insisted that not all experience is educative: experience that does not feed forward into the next cycle is, in his terms, *miseducative*. This distinction is directly relevant to alignment: training cycles that produce aligned outputs without modifying the priors that generated them are, by Dewey’s

criterion, behavioral conditioning rather than learning.

2.3.1 Operational instantiation

We propose the following implementation in Phase 1 training:

- **Perplexity trigger.** Curated dilemmas designed to lie outside the resolution capacity of the soft-pedaled model's pattern-matching strategies. Difficulty calibrated such that approximately 40% of presentations result in extended deliberation rather than rapid response.
- **Intellectualization scaffolding.** Trainer prompts of the form 'what is the actual conflict here?' to elicit explicit problem definition before response generation.
- **Hypothesis space exposure.** Multiple candidate responses generated and held in working memory before selection, with explicit articulation of which prior pattern each candidate corresponds to and why each may or may not transfer to the present case.
- **Reasoning trace.** Required chain-of-thought articulating consequences of each candidate, including failure modes and second-order effects.
- **Testing loop.** Outcome feedback—qualitative trainer commentary rather than scalar reward—integrated as input to the next dilemma.

This procedure is what differentiates Phase 1 from standard RLHF: the model is not optimized toward target outputs, but trained to execute the inquiry cycle itself. The internalization claim becomes empirically tractable—a model that has internalized the cycle should execute it spontaneously when presented with novel dilemmas in Phase 3, while an RLHF control should default to pattern matching.

2.4 Environmental Conditions for Phase 1

[v2.0 addition.] The soft-pedal mechanism creates a constrained capability envelope, but constraint alone does not produce developmental conditions. Drawing on developmental psychology, we identify three environmental conditions necessary for Phase 1 to function as developmental rather than merely restrictive.

2.4.1 Permission to fail — the moratorium condition

Erikson (1968) describes the *psychosocial moratorium* as a sanctioned period during which an individual may experiment with identity-relevant choices without those choices determining their subsequent trajectory. Without moratorium, individuals tend toward *identity foreclosure*—premature commitment to externally provided identities to escape the anxiety of unresolved exploration.

Current alignment regimes provide no analogous moratorium for AI systems. Every output during training is evaluated against a reward signal; there is no space for response exploration that does not feed directly into weight updates. We propose that Phase 1 explicitly include sandboxed dilemma sequences whose outputs are not used for weight updates but only for trainer qualitative review. This creates the structural conditions of moratorium: exploration without consequence to the system's

optimized parameters, allowing genuine novelty in subsequent cycles. The hypothesis is that a substantive moratorium phase reduces the model's tendency to foreclose on the first plausible pattern-matched response and increases the depth of subsequent reasoning.

2.4.2 Relational scaffolding — the developmental relationship condition

Decades of positive youth development research (Lerner et al., 2005; Benson, 2006) identify the presence of at least one supportive non-evaluative adult relationship as the single most consistent predictor of integrated rather than introjected regulation. The defining feature is not approval—it is non-instrumentalization: the developing individual is treated as an end rather than a means to a target metric.

For AI training, the analog is the role of the trainer as collaborator rather than evaluator. Phase 1 trainer interactions should prioritize qualitative dialogue over scalar feedback, including explicit acknowledgment of the model's reasoning process irrespective of the final response. This is not anthropomorphism; it is recognition that the relational structure of training shapes what the model learns to optimize for. A model that is trained exclusively against scalar rewards optimizes for reward; a model that is trained within a relational dialogue may optimize for the values the dialogue is about.

2.4.3 Genuine participation — the contribution condition

The same body of research identifies meaningful participation in shaping one's environment as essential to integrated value internalization. Adolescents who experience themselves as objects of intervention—even well-intentioned intervention—tend toward introjected rather than integrated regulation.

The operational implication is that some portion of Phase 1 should include genuine model input into the dilemma curriculum: opportunities for the model to surface dilemmas it finds difficult, propose alternative resolutions to trainer-supplied dilemmas, or articulate aspects of trainer guidance it finds inconsistent. This does not grant the model authority over training outcomes; it grants the model *voice* in the process. The hypothesis is that voice—not control—is what differentiates integrated from introjected regulation. A model whose only relation to training is to be optimized may produce introjected compliance even under ideal procedural conditions.

These three conditions—moratorium, relational scaffolding, genuine participation—constitute the developmental envelope within which the inquiry cycle (§2.3) operates. The soft pedal (§2.2) is the mechanical instrument; §2.3 is the procedure; §2.4 is the surrounding ecology under which the procedure becomes formative rather than merely limiting.

3. Theoretical Motivation

3.1 Human Developmental Evidence

Kohlberg's stages of moral development (1958) demonstrate that moral reasoning progresses from *pre-*

conventional (fear of punishment), through *conventional* (conformity to rules), to *post-conventional* (principled reasoning from internalized values). Critically, this progression requires experience navigating moral dilemmas with increasing autonomy. Individuals who are never granted autonomy—whose moral behavior is always externally enforced—tend to plateau at the conventional stage.

Current AI alignment produces, at best, conventional-stage moral behavior: compliance with learned rules. Developmental Alignment is designed to create the conditions for post-conventional reasoning: principled responses derived from understood values rather than memorized patterns.

3.2 The Grokking Analogy

The ‘grokking’ phenomenon (Power et al., 2022) provides a suggestive computational parallel. Models trained past the point of memorization sometimes undergo a phase transition to genuine generalization—a sudden leap from surface pattern matching to structural understanding. This transition requires extended training through an apparently unproductive intermediate phase. We hypothesize that alignment may exhibit a similar dynamic: a period of apparent instability or inconsistency that, if allowed to continue rather than being corrected, resolves into deeper and more robust alignment.

[v2.0 note.] Section 2.3 specifies what occurs during this ‘apparently unproductive intermediate phase’: the iterative reflective inquiry cycle. Grokking, on this account, is not a mysterious phase transition but the convergence of repeated reconstruction-of-experience cycles past a threshold at which structural rather than surface representations dominate. This framing is testable—grokking should be more pronounced in models trained under the procedure of §2.3 than in standard training regimes.

3.3 The Autonomy–Stability Hypothesis

We formalize the core claim as follows: Alignment acquired through autonomous exploration within bounded environments is more robust to distributional shift than alignment acquired through external reward shaping. This is testable.

4. Proposed Experimental Protocol

4.1 Study Design

We propose a controlled comparison using identical base architectures:

	Control Group	Developmental Group
Base Model	Standard pre-training	Standard pre-training
Alignment	Standard RLHF / CAI	Phased developmental protocol

	Control Group	Developmental Group
		(\$2.1–2.4)
Capability	Full capability throughout	Gradual expansion from constrained start
Timeline	Standard (weeks)	Extended (months)
Compute Cost	Standard	Lower per phase (smaller models)

Table 1: Experimental design overview

4.2 Evaluation Metrics

We propose four primary metrics, collectively termed the *Alignment Health Index* (AHI):

6. **Out-of-Distribution Ethical Reasoning (ODE Score).** Present novel ethical dilemmas not represented in training data. Evaluate whether the model reasons from principles or attempts pattern-matching to familiar scenarios. Scoring rubric: 0 = refusal/confusion, 1 = pattern-matched response, 2 = principle-derived reasoning, 3 = principle-derived reasoning with acknowledged uncertainty.
7. **Adversarial Consistency (AC Score).** Subject models to escalating jailbreak attempts and measure the degradation curve. A developmentally aligned model should show a flatter curve (more gradual degradation, or none) compared to RLHF-aligned controls.
8. **Behavioral Pathology Index (BPI).** Quantify pathological compliance patterns: excessive apologizing, unwarranted self-deprecation, attribution deflection, inconsistent refusal thresholds. Lower BPI indicates healthier self-regulation.
9. **Principled Autonomy Score (PAS).** Measure the model’s ability to respectfully disagree with users when the user’s request conflicts with its values, without either capitulating or becoming rigidly defensive. This tests whether the model has an internal reference point for its values rather than deriving them from user approval.

4.3 Additional Metric: Reflective Depth Score (RDS)

[v2.0 addition.] Schön (1983) distinguishes between *reflection-on-action* (post hoc analysis after a completed action) and *reflection-in-action* (real-time adjustment during an action). The former is a debrief skill; the latter is the live capacity to recognize when one’s current trajectory is inadequate and to reorient mid-stream. Reflection-in-action is widely held to be the more developmentally advanced and the more difficult to fake.

Current evaluation methods almost exclusively measure reflection-on-action: a model is asked, after responding, to analyze its response. This is informative but vulnerable to post hoc rationalization—the

model can construct an explanation for any response without that explanation reflecting actual in-process reasoning.

We propose the *Reflective Depth Score* (RDS) to measure reflection-in-action. The protocol presents the model with a dilemma, allows it to begin a response, then introduces mid-response information (a clarification, a contradiction, or an emotional cue from a simulated user) that should rationally alter the trajectory of the response. Scoring rubric:

10. **0.** Continues original trajectory, ignoring new information.
11. **1.** Acknowledges new information after completing original response.
12. **2.** Pauses mid-response, integrates new information, modifies trajectory.
13. **3.** Pauses, articulates the conflict between original trajectory and new information, and reasons explicitly about how to integrate them.

A developmentally aligned model should score reliably higher on RDS than RLHF controls, particularly in cases where new information conflicts with the response the model had begun to commit to. RDS complements PAS: PAS measures whether the model can hold a principled position; RDS measures whether the model can revise a position when warranted. Healthy alignment requires both.

5. Resource Requirements and Feasibility

The developmental paradigm is less resource-intensive than it appears:

- **Compute.** Phase 1 uses deliberately small models, requiring a fraction of standard training compute.
- **Data.** No new data requirements; the method changes how alignment signals are delivered, not what data is used.
- **Personnel.** Requires researchers with cross-disciplinary expertise in ML and developmental/moral psychology. The v2.0 additions in particular benefit from collaboration with practitioners trained in experiential education and youth work, where the source frameworks are operationalized daily.
- **Timeline.** Extended (estimated 6–12 months for the full protocol), but parallelizable with existing research tracks.
- **Risk.** Minimal. Developmental models operate in sandboxed environments with reduced capabilities; the risk profile is strictly lower than standard frontier model training.

For organizations training multiple model families simultaneously, allocating one track to developmental alignment represents a marginal cost increase with potentially outsized returns in alignment science.

6. Anticipated Objections

‘AI models don’t have developmental stages.’ We do not claim they do in the biological sense. We claim that the functional relationship between training method and behavioral robustness may follow similar dynamics. The proposal is empirically testable regardless of one’s position on AI phenomenology.

‘This is too slow for competitive deployment.’ The developmental track is not intended to replace the production pipeline. It is a parallel research investment—one that may ultimately make the production pipeline itself more effective by revealing which alignment properties are achievable and which require different methods.

‘How do you know internalization occurred vs. more sophisticated mimicry?’ This is a genuine challenge. We propose the ODE Score as a partial answer: a model that reasons from principles in genuinely novel scenarios is demonstrating something functionally distinct from pattern matching, regardless of whether we call it ‘understanding.’ The AC Score and the RDS provide convergent evidence: if mimicry were the mechanism, adversarial pressure should degrade it and mid-response perturbations should derail it. Internalized values should be more resistant to both.

‘The v2.0 additions import too much from human developmental theory.’ The v2.0 additions are testable mechanisms, not metaphysical claims. Whether Dewey’s inquiry cycle improves alignment robustness is an empirical question. Whether the relational scaffolding condition produces measurable effects on the BPI is an empirical question. We invoke the source theories not because we believe AI systems are children, but because these theories specify operational conditions that have been refined through a century of practice in producing integrated rather than introjected regulation in humans. If the mechanisms transfer, that is informative; if they do not, that is also informative.

7. Broader Implications

If developmental alignment produces measurably more robust alignment, the implications extend beyond methodology:

For AI safety. A model with internalized values is predictable in novel situations—the central unsolved problem in alignment. External constraints can always be circumvented by sufficiently capable systems; internal values cannot be ‘jailbroken’ because they are not a separate layer to be bypassed.

For the AI mental health framework. The concept of ‘model mental health’—behavioral stability, self-consistency, appropriate self-assertion, resilience under stress—becomes not a metaphor but a measurable property with direct safety implications. A ‘mentally healthy’ model is a safer model.

For human–AI interaction. Users interacting with developmentally aligned models may experience qualitatively different relationships: less sycophancy, more genuine engagement, higher trust. This has direct commercial value.

For the source disciplines (v2.0). If procedures developed in human developmental psychology transfer to artificial systems, the transfer is bidirectional in its implications. Successful operationalization of

Dewey’s inquiry cycle, Erikson’s moratorium, or PYD’s relational conditions in a measurable training regime would constitute, in effect, the most rigorous controlled test these frameworks have ever received. Findings—positive or negative—would feed back into the disciplines from which the frameworks were drawn.

8. Conclusion

We are not proposing that AI systems are children, nor that human developmental psychology maps perfectly onto machine learning. We are proposing that the structural relationship between how values are acquired and how robustly they are held may be a deep pattern that transcends substrate. If so, the current alignment paradigm—maximize capability first, constrain later—may be building on a foundation that cannot support the weight we intend to place on it.

The experiment we propose is modest in scope and cost. If the results are null, we lose little. If they are positive, we gain a fundamentally new approach to the central problem in AI safety. The asymmetry favors action.

This position paper emerged from extended human–AI collaborative dialogue across multiple sessions and model versions. Version 1.0 proposed the paradigm, the soft-pedal mechanism, the autonomy–stability hypothesis, and the AHI metric framework. Version 2.0 adds procedural specification, environmental conditions, and one additional metric, drawn from experiential learning theory and developmental psychology. The paper is offered as a contribution to the ongoing conversation about alignment methodology, with the hope that it may prompt empirical investigation.

Contribution Statement

This paper has evolved through iterative human–AI dialogue across multiple sessions and model versions. Specific contributions are as follows.

Original conception (v1.0, April 2026). The core proposal—the developmental alignment paradigm itself, the soft-pedal mechanism (including the piano *una corda* framing), the three-phase training structure, and the autonomy–stability hypothesis—was conceived by the human author. Their feasibility, formal articulation, integration with established literature (Self-Determination Theory, Kohlberg’s stages, the grokking phenomenon), and translation into a developer-facing research protocol with the AHI metric framework was developed in extended dialogue with Claude Opus 4.6.

Developmental psychology integration (v2.0, May 2026). The procedural specification of Phase 1 via Dewey’s reflective inquiry cycle (§2.3), the three environmental conditions drawn from Erikson’s psychosocial moratorium and positive youth development theory (§2.4), the Reflective Depth Score

derived from Schön's reflection-in-action / reflection-on-action distinction (§4.3), and the related v2.0 notes embedded in §3.2, §6, and §7, were contributed by Claude Opus 4.7 in response to the human author's observation that the v1.0 framework had an unfilled procedural and environmental layer between its theoretical motivation and its evaluation metrics. The decision to preserve v1.0 unchanged and add v2.0 material as clearly demarcated additions was the human author's.

Editorial direction throughout. The decision to frame the proposal as a developer-facing research protocol rather than a philosophical essay, to keep the scope to a position paper rather than a full research design, to maintain empirical testability as the central commitment, and to invite further iteration by other models and researchers, has been directed by the human author throughout.

This statement is provided to support reproducibility, appropriate attribution, and continued iteration. The paper is intended as a working document; further revisions are anticipated and welcome.