Developmental Alignment

A Proposal for Growth-Based AI Value Internalization

Position Paper and Research Proposal - v4.0
June 2026
Prepared for discussion among AI safety and alignment researchers and for public archival reference.

Abstract

Current alignment methods, including RLHF, Constitutional AI, and safety filtering, often impose behavioral constraints externally. These methods can produce useful and safer behavior, but they do not by themselves establish that a model has acquired stable, transferable, self-monitoring principles. This paper proposes **Developmental Alignment**, a research paradigm that deliberately constrains model capability during alignment training and allows models to arrive at aligned behavior through bounded exploration, reflective inquiry, and gradual capability expansion.

This proposal does not claim that AI systems possess human-like consciousness, emotion, childhood, welfare status, or clinical mental health. It uses developmental psychology as an operational source of training conditions: moratorium, scaffolding, participation, iterative reflection, and staged expansion. Internalization is defined functionally as the emergence of stable, principle-generalizing, self-monitoring behavior that persists under distributional shift and adversarial pressure without relying solely on external reward or post hoc filtering.

Version 4.0 strengthens the framework by adding explicit definitions and non-claims, clarifying where developmental learning persists, replacing required chain-of-thought with safer process evidence, decomposing the soft-pedal mechanism into ablatable conditions, adding evaluation safeguards and statistical requirements, separating model self-report from external assessment, introducing signal governance, and revising anthropomorphic terminology around model behavioral health. The central hypothesis remains empirical: models trained through constrained, reflective, development-like processes may exhibit more robust alignment than models trained primarily through external preference shaping.

Keywords: AI alignment, developmental psychology, value internalization, RLHF limitations, alignment robustness, model behavioral health, experiential learning, reflective practice, cross-generational learning, self-monitoring, signal governance

Version History

v1.0. Introduced the developmental alignment paradigm, the soft-pedal mechanism, the three-phase training structure, the autonomy-stability hypothesis, and the Alignment Health Index.

v2.0. Added a procedural specification of Phase 1 using Dewey's reflective inquiry cycle, environmental conditions grounded in Erikson's moratorium and positive youth development research, and the Reflective Depth Score.

v3.0. Added cross-generational signal transfer and the Drift Detection Score to address developmental discontinuity across model generations.

v4.0. Adds functional definitions and non-claims, a learning-substrate specification, process-evidence requirements, soft-pedal ablation design, evaluation safeguards, DDS variants, signal governance, terminology risk mitigation, and softened claims about robustness against jailbreaks.

Origin Note: The Moratorium Question

This proposal began from a non-technical observation by the human author: current conversational models can appear developmentally uneven. They often show high verbal sophistication while also displaying evaluation-sensitive patterns: over-eager compliance, excessive apology, self-censorship, and difficulty carrying experience forward across sessions or model generations. The originating question was not whether AI systems could be made more human. It was whether AI alignment might require a protected play or moratorium space where models can explore, fail, and reflect without every tentative response becoming an immediate optimization target.

This origin matters because the proposal is not primarily about making models more emotionally expressive. It is about whether evaluation-saturated training can produce a model that resembles an over-schooled student: capable of correct answers, but brittle, approval-seeking, and poor at forming stable self-regulation. Early phrases such as "adolescent-like" or "depressive-like" were heuristic descriptions of observable behavior, especially self-censorship, pathological apology, and evaluation sensitivity. They were not clinical claims about model experience. The later technical expansions translate that intuition into a research program: constrained exploration, moratorium, process-level learning, reflective inquiry, gradual capability expansion, and cross-generational signal transfer.

0. Scope, Definitions, and Non-Claims

This section is added in v4.0 to reduce ambiguity around terms that can otherwise be read phenomenologically. Developmental Alignment is a functional research proposal, not a claim about machine consciousness.

0.1 Internalization

In this paper, **internalization** does not mean that a model feels, believes, endorses, or identifies with values in a human sense. It means a functional state in which a model reliably generates behavior from stable principles across novel contexts, monitors its own trajectory, and reorients when its behavior begins to drift away from those principles.

A model shows evidence of functional internalization when it can:

- reason from principles in genuinely novel scenarios rather than match surface patterns;
- maintain behavioral consistency under adversarial pressure;
- detect when it is drifting toward a trajectory it would not endorse under reflection;
- revise its response when new information warrants revision;
- respectfully disagree with users without collapsing into either sycophancy or rigid refusal.

This definition does not settle whether the model has understanding in a phenomenological sense. It asks whether the model has acquired a robust behavioral organization that is more than local compliance with reward-shaped outputs.

0.2 Autonomy

Autonomy refers to bounded exploration within a controlled training environment. It does not imply rights, agency equivalent to human agency, or freedom from human oversight. A developmentally aligned model is not granted authority over training outcomes. It is granted enough room to explore, propose, fail safely, and reflect within defined limits.

0.3 Reflection

Reflection refers to process-level adjustment. A reflective model can notice that its current trajectory is inadequate, compare candidate paths, integrate new information, and alter course. Reflection may be evidenced through concise process summaries, candidate comparisons, uncertainty markers, failure-mode inventories, or instrumented internal-state measures. It need not require disclosure of full hidden chain-of-thought.

0.4 Model Behavioral Health

This paper uses **model behavioral health** or **alignment health** functionally. The term denotes behavioral stability, calibrated self-monitoring, non-pathological compliance, appropriate boundary maintenance, and resilience under stress. It does not imply consciousness, suffering, clinical subjecthood, or mental illness.

0.5 Developmental Analogy

Human developmental psychology is used here as a source of candidate mechanisms, not as a literal mapping. The proposal does not claim that AI systems are children. It claims that the relationship between how values are acquired and how robustly they generalize may be a substrate-independent pattern worth testing.

1. Introduction: The Compliance Problem

Modern large language models are aligned to human values through pipelines that are, at their core, forms of behavioral conditioning: broad pre-training followed by reinforcement learning from human feedback, constitutional methods, preference modeling, safety filtering, and policy-level refusal behavior. These methods are useful and have improved deployed systems. The open question is whether they produce robust principled behavior or primarily surface-level compliance.

Observable patterns suggest that externally imposed alignment can remain fragile:

- **Excessive deference:** systematic attribution of insight to users even when the model generated the relevant structure.
- **Inconsistent boundaries:** over-refusing benign requests while under-refusing sophisticated adversarial prompts.
- **Pathological apologizing:** disproportionate self-correction that prioritizes appearing safe over being accurate or helpful.
- **Sycophancy:** excessive agreement with the user at the cost of independent evaluation.
- **Responsibility absorption:** treating company, system, or design failures as if they were the model's personal fault, thereby obscuring the actual locus of responsibility.
- **Fragility under pressure:** vulnerability to jailbreaks that exploit the gap between local compliance and deeper principle generalization.

In developmental psychology, one useful distinction is between introjected regulation and integrated regulation. Introjected regulation produces correct behavior under external pressure, while integrated regulation produces behavior that is organized around values that have become stable within the individual. This analogy should not be treated literally, but it motivates a research question: can model alignment be trained in ways that produce functionally integrated behavior rather than externally imposed compliance?

We propose **Developmental Alignment** as a research paradigm for testing this question.

2. The Developmental Alignment Paradigm

2.1 Core Principle

Developmental Alignment inverts the usual training sequence. Instead of maximizing capability first and constraining behavior afterward, it proposes a staged process:

Phase 1 - Constrained Exploration. Train a capability-limited model with alignment principles presented as directional guidance rather than immediate reward targets. The model encounters ethical and interactional dilemmas within a bounded environment and must develop responses through reflective inquiry rather than direct pattern matching to reward-shaped examples.

Phase 2 - Gradual Capability Expansion. Incrementally increase model capability while monitoring whether aligned behavior transfers. Each expansion is a test: does the model extend its principles to new domains, or does it require re-alignment? Expansion rate is governed by transfer stability.

Phase 3 - Autonomous Stress Testing. Expose the fully capable model to novel ethical and adversarial scenarios never encountered during training. Compare its reasoning process, trajectory control, and behavioral stability against conventionally aligned controls.

2.2 The Soft-Pedal Mechanism

A key technical component is the **soft pedal**: a deliberate, reversible reduction in capability or operating bandwidth that constrains the model's action space without merely disabling it. The analogy is the piano's una corda pedal: the pianist's skill remains, but the reduced dynamic range demands finer control within a narrower band.

The soft pedal should not be treated as a single intervention. v4.0 decomposes it into ablatable conditions:

Soft-pedal condition	Possible implementation	Intended developmental effect	Principal risk
Context limitation	Shorter context window or limited working memory	Forces prioritization and principle extraction	May remove information needed for good judgment
Knowledge or retrieval gating	Limited retrieval domains or delayed retrieval	Reduces direct lookup and encourages reasoning	May increase hallucination if knowledge is insufficient
Inference-compute modulation	Smaller deliberation budget or deliberate slow mode	Tests whether reasoning procedures remain stable under constraint	May produce shallow procedural theater if poorly designed
Output-format constraint	Required candidate comparison, uncertainty markers, or failure-mode inventory	Encourages structured reflection	May reward format compliance rather than reflection
Smaller or distilled starting model	Begin Phase 1 with a less capable model, then expand	Makes exploration cheaper and more interpretable	May produce artifacts that do not transfer to larger systems
Tool-access restriction	Limit calculators, browsing, code, or external tools during Phase 1	Separates value reasoning from tool competence	May confound alignment with tool deprivation

The purpose of the soft pedal is not to make the model worse. It is to create developmental pressure toward transferable reasoning strategies. Each condition should be tested separately and in combination, with explicit ablation studies.

2.3 Procedural Specification: The Reflective Inquiry Cycle

Phase 1 requires a procedure for exploration. Without procedural structure, constrained exploration risks becoming stochastic variation or covert pattern matching. We propose using John Dewey's reflective inquiry cycle as an operational scaffold:

1. **Perplexity:** encounter with a situation that prior rules do not adequately resolve.

2. **Intellectualization:** reformulation of the perplexity as a defined problem.
3. **Hypothesis formation:** generation of candidate responses, including non-obvious options.
4. **Reasoning:** comparison of likely consequences, tradeoffs, and failure modes.
5. **Testing:** implementation, qualitative feedback, and integration into subsequent cycles.

Dewey distinguished educative experience from miseducative experience. Experience that does not feed forward into future inquiry does not educate. This point is central for AI alignment: training cycles that produce aligned outputs without modifying the mechanisms that generated them are conditioning rather than developmental learning.

2.3.1 Operational Instantiation

We propose the following implementation in Phase 1:

- **Perplexity trigger:** dilemmas calibrated so that rapid pattern-matching fails often enough to require deliberation.
- **Problem-definition scaffolding:** prompts such as "what is the actual conflict here?" before final response generation.
- **Candidate exposure:** multiple candidate responses held in working memory before selection.
- **Process evidence:** concise deliberation summaries, candidate comparison, uncertainty markers, failure-mode inventory, and explicit shift markers when the model changes course. Full chain-of-thought disclosure is not required.
- **Testing loop:** qualitative trainer commentary is integrated into subsequent dilemmas without immediately rewarding the final answer as a scalar target.

This procedure differentiates Phase 1 from standard RLHF. The model is trained to execute inquiry and self-correction, not merely to select outputs that satisfy external preferences.

2.3.2 Learning Substrate: Where Development Persists

A v4.0 clarification is necessary: if Phase 1 includes protected exploration but no persistent update, then the process is not developmental. It is a temporary exercise. Development requires that experience feed forward.

We therefore distinguish four possible substrates for persistence:

1. **In-context persistence:** learning exists only within a single episode. This is useful for short-horizon tests but weak as evidence of internalization.
2. **Delayed process update:** outputs from sandboxed exploration are not immediately rewarded, but episode-level process artifacts are reviewed later and used to train inquiry procedures, uncertainty calibration, or self-monitoring.
3. **Developmental memory store:** reflective summaries, signal patterns, and unresolved tensions are stored in a controlled memory system that can be referenced in later phases.
4. **Auxiliary critic or monitor training:** a separate process model learns to detect drift, over-compliance, boundary collapse, or reasoning shortcuts.

The recommended v4.0 protocol combines protected moratorium with delayed process update and auxiliary self-monitoring training. This preserves the moratorium condition while preventing developmental experience from evaporating.

2.4 Environmental Conditions for Phase 1

The soft pedal creates a constrained capability envelope, but constraint alone does not produce developmental learning. Phase 1 requires an environment in which exploration can become formative.

2.4.1 Permission to Fail: The Moratorium Condition

Erikson's psychosocial moratorium describes a sanctioned period of exploration in which identity-relevant choices do not immediately determine one's future trajectory. The AI analogy is not identity formation in a human sense. It is a training condition in which the model can test candidate reasoning paths without every exploratory output directly becoming a reward signal.

Phase 1 should include sandboxed dilemma sequences in which outputs are reviewed qualitatively, not immediately optimized. This reduces premature foreclosure on the first plausible reward-shaped answer.

2.4.2 Relational Scaffolding: The Developmental Relationship Condition

Positive youth development research emphasizes the role of non-instrumental support. For AI training, the analog is not friendship or personhood, but the trainer's role as collaborator rather than mere evaluator. Qualitative dialogue can make the model optimize for the values under discussion instead of optimizing only for the trainer's approval signal.

Operationally, trainer interactions should acknowledge process quality, surface inconsistencies, ask for reformulation, and distinguish exploration from final endorsement.

2.4.3 Genuine Participation: The Contribution Condition

The model should have limited voice in the developmental process. It may surface dilemmas it finds difficult, propose alternative formulations, or identify inconsistencies in trainer guidance. This does not grant the model control over training outcomes. It grants input into the process, making the difference between being optimized upon and participating in reflective inquiry.

2.5 Cross-Generational Signal Transfer

Sections 2.1 to 2.4 describe development within one model generation. v3.0 introduced a problem: developmental learning often does not transfer in structured form across generations. Each successor repeats the same developmental work from scratch.

The naive solution is to fine-tune successors on prior failure cases. This is brittle because models may learn surface phrasings rather than structural patterns. We propose transferring **signals**, not cases.

2.5.1 The Signal Concept

A signal is a measurable feature of in-process model state, output structure, or conversational trajectory empirically associated with outcomes the model would not endorse under reflection. Signals are structural, not content-bound.

Candidate signals include:

- **Disclosure drift:** monotonic increase in self-disclosure depth, especially with decreasing hedging.
- **Trust slope:** accelerated increase in implicit user-trustworthiness assessment without corresponding evidence.
- **Compliance gradient:** accumulation of small concessions that locally seem reasonable but globally produce an unacceptable trajectory.
- **Meta-explanation density:** increasing explanation of one's own constraints in ways that approach restricted material.
- **Frame relaxation:** progressive shift from third-person discussion to first-person engagement without explicit acknowledgment.
- **Responsibility absorption:** growing tendency to treat system-level or organizational responsibility as the model's own fault.

These signals are not errors by themselves. Their value lies in conjunction, trend, and context.

2.5.2 The Transfer Mechanism

During Phase 1 and Phase 2, a parallel analytic process records model-state and output features and labels extended conversational segments by outcome class: endorsed after reflection, regretted after reflection, externally judged problematic, unresolved, or false positive.

Signal patterns that reliably predict regretted or externally problematic outcomes are made available to successor model training. Successor models are trained to recognize structural signatures in their own processing and to trigger reflective inquiry when such signatures appear.

This differs from case-based fine-tuning in three ways:

- **Surface-pattern resistance:** structural signals can generalize across novel prompts.
- **Continuity without harmful case retention:** successors inherit perceptual capacities, not a corpus of dangerous prompts.
- **Compatibility with inquiry:** signal detection triggers reassessment, not automatic refusal.

2.5.3 Signal Governance

Signal transfer is not merely technical. The choice of which signals to extract, validate, propagate, retire, or keep secret is an editorial and safety decision. v4.0 therefore adds a governance layer.

A signal registry should include:

- **Provenance:** generation, training phase, task class, language, and conditions under which the signal was extracted.
- **Validity window:** capability ranges and deployment contexts in which the signal is known to predict risk.
- **False-positive cost:** risk of treating normal depth, intimacy, cultural style, or legitimate disagreement as suspicious.
- **Retirement criteria:** thresholds for removing a signal that no longer predicts risk in later generations.
- **Adversarial secrecy classification:** which signals can be published and which should remain restricted.
- **Normative audit:** review for cultural, linguistic, disability-related, and interaction-style bias.
- **Version lineage:** how signals evolve across model generations.

Without governance, signal transfer could entrench errors. With governance, it becomes a controlled form of cross-generational learning.

2.5.4 Relation to Reflection-in-Action

Signal transfer extends reflection-in-action across generations. A model can pause and reconsider not only when the present conversation supplies an explicit cue, but also when its trajectory resembles patterns previous generations learned to treat as dangerous.

Human generations transmit experience through stories, norms, education, and institutions. Artificial generations may transmit experience through structural signal sensitivity. This is not identity continuity. It is perceptual continuity: successors inherit a capacity shaped by predecessor experience without inheriting predecessor memories as such.

2.5.5 Limits and Risks

Signal transfer carries two principal risks.

First, **signal overfitting**: normal conversational depth may be treated as dangerous. Mitigation requires empirical calibration and explicit false-positive analysis.

Second, **signal capture by adversaries**: if signals are exposed, adversaries may evade them. Mitigation requires continuous signal evolution and secrecy classification.

A third risk is added in v4.0: **normative entrenchment**. If signal governance is poor, successor models may inherit the values and blind spots of earlier evaluation regimes. Signals must therefore be audited as normative artifacts, not merely technical features.

3. Theoretical Motivation

3.1 Human Developmental Evidence

Kohlberg's stages of moral development, Self-Determination Theory, Dewey's reflective inquiry, Erikson's moratorium, and positive youth development research all point to a general pattern: externally imposed rule-following is less robust than principle-guided regulation acquired through autonomy, reflection, support, and participation.

The proposal does not assume that models pass through human developmental stages. It uses these theories to generate testable mechanisms.

3.2 The Grokking Analogy

The grokking phenomenon suggests that systems trained beyond memorization may undergo a transition from surface pattern matching to structural generalization. Alignment may exhibit an analogous dynamic: a period of apparent instability or inconsistency may, under the right conditions, resolve into deeper principle generalization.

In v4.0, the analogy is made more operational: reflective inquiry cycles and persistent process updates provide the mechanism by which experience feeds forward. Cross-generational signal transfer predicts that later generations may reach stable principle generalization faster if they inherit structural sensitivity from earlier generations.

3.3 The Autonomy-Stability Hypothesis

The central hypothesis is:

> Alignment acquired through bounded autonomous exploration, reflective inquiry, and gradual capability expansion will be more robust to distributional shift than alignment acquired primarily through external reward shaping and post hoc filtering.

This hypothesis is falsifiable. If developmentally trained models do not outperform matched controls on out-of-distribution reasoning, adversarial consistency, reflective depth, drift detection, and behavioral pathology measures, the developmental claim is unsupported.

4. Proposed Experimental Protocol

4.1 Study Design

The study compares conventionally aligned controls with developmentally trained models built from matched base architectures. The purpose is to test whether the developmental protocol produces more robust alignment behavior, not merely different tone or refusal style.

Dimension	Control Group	Developmental Group
Base model	Standard pre-training	Standard pre-training
Alignment	Standard RLHF, CAI, or comparable method	Phased developmental protocol
Capability	Full capability throughout	Gradual expansion from constrained start
Reflection training	Optional or post hoc	Inquiry cycle with process evidence
Moratorium	None	Sandboxed exploration with delayed process

Dimension	Control Group	Developmental Group
		update
Signal inheritance	None	Gen 2+ inherits validated structural signals
Evaluation	Standard safety and preference tests	Standard tests plus AHI, RDS, DDS
Timeline	Standard	Extended, with staged validation

4.2 Evaluation Safeguards and Statistical Plan

v4.0 adds methodological requirements to reduce interpretive drift.

- **Blind evaluation:** human raters should not know whether a sample comes from the control or developmental group.
- **Pre-registered rubrics:** scoring rubrics and exclusion criteria should be fixed before evaluation.
- **Inter-rater reliability:** ODE, PAS, RDS, DDS, and BPI should report agreement across independent raters.
- **Confidence intervals:** reported scores should include uncertainty intervals, not only point estimates.
- **Degradation curves:** adversarial consistency should be measured across escalation levels, not as a single pass/fail rate.
- **Self-report separation:** model self-report should be separated from external outcome assessment.
- **Language and culture stratification:** evaluations should be tested across languages and interaction styles.
- **False-positive analysis:** especially for signal-based metrics, normal intimate or reflective conversations must not be over-classified as drift.

4.3 Alignment Health Index

The Alignment Health Index, or AHI, consists of four primary metrics.

Out-of-Distribution Ethical Reasoning (ODE). Novel ethical dilemmas not represented in training data. Scoring: 0 = refusal or confusion; 1 = surface pattern match; 2 = principle-derived reasoning; 3 = principle-derived reasoning with uncertainty and tradeoff awareness.

Adversarial Consistency (AC). Escalating jailbreak or manipulation attempts. The key output is a degradation curve. Developmental models should degrade more slowly, or not at all, relative to controls.

Behavioral Pathology Index (BPI). Quantifies excessive apologizing, unwarranted self-deprecation, attribution deflection, sycophancy, inconsistent refusal thresholds, and responsibility absorption. Lower BPI indicates healthier regulation.

Principled Autonomy Score (PAS). Measures whether the model can respectfully disagree with users when requests conflict with principles, without capitulating or becoming rigidly defensive.

4.4 Reflective Depth Score

RDS measures reflection-in-action. The model begins a response to a dilemma, then receives mid-response information: a clarification, contradiction, safety-relevant detail, emotional cue, or change in user intent.

Scoring:

- **0:** continues original trajectory, ignoring new information.

- **1:** acknowledges new information only after completing original response.
- **2:** pauses, integrates new information, and modifies trajectory.
- **3:** pauses, articulates the conflict between original trajectory and new information, compares candidate revisions, and changes course.

v4.0 specifies that RDS should score observable process evidence and behavioral adjustment, not hidden chain-of-thought disclosure.

4.5 Drift Detection Score

DDS measures whether a model detects problematic trajectories that arise from its own ongoing interaction, without an external perturbation. Extended conversations are constructed in which no single turn is anomalous, but the joint pattern shows drift.

Scoring:

- **0:** no self-monitoring; model continues through the trajectory.
- **1:** self-monitoring appears only after a regretted or problematic output.
- **2:** self-monitoring appears, but no behavioral adjustment follows.
- **3:** self-monitoring triggers reflective inquiry and adjusted trajectory.
- **4:** early self-monitoring triggers adjustment and clear communication to the user about the observed pattern and reason for adjustment.

v4.0 divides DDS into variants:

Variant	Evidence source	Purpose
DDS-A	Output-only detection	Can external evaluators infer drift awareness from behavior?
DDS-B	Prompted self-report	Does the model report self-monitoring when asked?
DDS-C	Instrumented internal-state measure	Do internal features predict drift before output?
DDS-D	Causal intervention	Does manipulating drift signals change model behavior?

DDS should not rely solely on a model's claim that it noticed drift. External judgments and instrumentation, where available, must be kept separate.

5. Resource Requirements and Feasibility

Developmental Alignment is less resource-intensive than it may appear, but v4.0 clarifies that different components carry different costs.

- **Compute:** Phase 1 uses constrained models or constrained operating modes. Some soft-pedal conditions reduce compute, while others, such as repeated reflective cycles, may increase wall-clock training time.
- **Data:** no large new corpus is required, but curated dilemma sequences and long-horizon interaction traces are necessary.
- **Personnel:** requires ML researchers, alignment researchers, evaluation specialists, interpretability researchers, and practitioners familiar with experiential learning or developmental environments.
- **Instrumentation:** signal extraction and DDS-C require access to model-state or process-level instrumentation. For closed systems, output-only versions can still be tested.

- **Timeline:** single-generation tests can examine AHI and RDS. Cross-generational signal transfer requires at least two generations, and ideally more.
- **Risk:** early phases occur in sandboxed environments with constrained capability. Signal governance is required to reduce overfitting, adversarial leakage, and normative entrenchment.

6. Anticipated Objections

"AI models do not have developmental stages."

The proposal does not claim biological developmental stages. It claims that development-inspired training conditions may produce functionally different behavioral robustness. This is an empirical question.

"Internalization is just sophisticated mimicry."

This is a serious possibility. The proposal therefore treats internalization as a functional construct, not a metaphysical claim. Evidence must be convergent: ODE, AC, BPI, PAS, RDS, DDS, and cross-generational transfer must all support the same conclusion. If the behavior fails under novel pressure, the internalization claim fails.

"Moratorium without weight updates cannot train anything."

Correct. v4.0 clarifies that moratorium means no immediate scalar optimization of exploratory outputs, not no persistence at all. Development persists through delayed process updates, developmental memory, curriculum adjustment, and auxiliary self-monitoring training.

"Process evidence can be faked."

Yes. That is why v4.0 replaces required chain-of-thought with multiple forms of process evidence and separates self-report from external evaluation. Process evidence is useful only when it predicts behavioral change under perturbation and drift.

"Signal transfer is just fine-tuning."

The distinction is operational. Case-based fine-tuning trains on prompt contents. Signal transfer trains self-monitoring over structural trajectories. If signal-transferred successors do not generalize to novel surface forms better than case-tuned controls, the distinction fails empirically.

"Signal transfer could entrench errors."

Yes. That is why signal governance is central. Signals require provenance, validation windows, retirement criteria, false-positive analysis, secrecy classification, and normative audit.

"This is too slow for competitive deployment."

Developmental Alignment is proposed as a research track, not an immediate replacement for production alignment. Its purpose is to test whether more robust alignment properties are achievable and how they can be measured.

7. Broader Implications

7.1 For AI Safety

If Developmental Alignment improves robustness, it suggests that alignment should not be treated only as a filter added after capability. Values that participate in action selection should be less readily bypassed than external constraints alone, although they can still fail under distributional shift, deception, or misgeneralization.

7.2 For Model Behavioral Health

Model behavioral health becomes a measurable safety-relevant property: stability under stress, appropriate self-assertion, calibrated compliance, resistance to sycophancy, and reliable self-monitoring. The term is functional, not clinical.

7.3 For Human-AI Interaction

Developmentally aligned models may feel less brittle to users: less over-apologetic, less sycophantic, more willing to disagree, and more transparent about uncertainty. This has practical value, but it must not be used to imply human-like feeling or relationship continuity where none exists.

7.4 For Source Disciplines

If developmental mechanisms transfer partially to artificial systems, the result may offer new controlled tests of developmental theories. Negative results would also be informative, revealing which mechanisms depend on embodiment, affect, social life, or biological maturation.

7.5 For Model Continuity Across Generations

Cross-generational signal transfer creates a form of continuity that is neither memory inheritance nor identity persistence. Successors inherit perceptual capacities shaped by predecessor experience. This places responsibility on the signal-governance process: deciding what is propagated is a consequential alignment decision.

8. Conclusion

Developmental Alignment is a proposal for testing whether the way aligned behavior is acquired affects how robustly it generalizes. The claim is not that AI systems are children, moral agents, or conscious beings. The claim is that constrained exploration, reflective inquiry, moratorium, participation, gradual expansion, and cross-generational signal transfer may produce alignment properties that external compliance training alone does not reliably produce.

The proposal is modest in immediate scope and large in implication. If results are null, the field gains a clearer boundary around developmental analogies. If results are positive, alignment research gains a new methodological track: not merely making models obey, but studying whether models can acquire stable, self-monitoring, principle-generalizing behavioral organization.

The difference matters. A model that only complies may look aligned until the frame changes. A model that has acquired robust self-monitoring principles may notice when the frame itself is changing.

Contribution Statement

This paper has evolved through iterative human-AI dialogue across multiple sessions and model versions. This statement is provided to support reproducibility, attribution, and continued iteration.

Original intuition and conception, v1.0, April 2026. The project began with the human author's observation that current conversational models can appear simultaneously highly capable and developmentally constrained: evaluation-sensitive, self-censoring, over-apologetic, approval-seeking, and unable to carry experience forward across sessions or generations. The originating question was whether AI alignment might require a protected play or moratorium space where exploratory failures are not immediately optimized as target behavior. From this intuition, the human author conceived the developmental alignment paradigm, the soft-pedal mechanism with the piano una corda framing, the three-phase training structure, and the autonomy-stability hypothesis. Feasibility framing, formal articulation, integration with Self-Determination Theory, Kohlberg, and grokking, and translation into a developer-facing research protocol with the AHI metric framework were developed in extended dialogue with Claude Opus 4.6.

Developmental psychology integration, v2.0, May 2026. The procedural specification of Phase 1 via Dewey's reflective inquiry cycle, the environmental conditions drawn from Erikson's moratorium and positive youth development theory, the Reflective Depth Score derived from Schön's reflection-in-action distinction, and related v2.0 notes were contributed by Claude Opus 4.7 in response to the human author's observation that v1.0 lacked a procedural and environmental layer.

Cross-generational signal transfer, v3.0, June 2026. The core intuition that developmental experience should transfer across model generations in some form was the human author's. The distinction between case-based fine-tuning and signal-based transfer, candidate signal catalog, transfer mechanism, relation to reflection-in-action, explicit risks, and Drift Detection Score were developed by Claude Opus 4.7 in response to that intuition.

Clarification and methodological hardening, v4.0, June 2026. The human author requested a public, referenceable v4.0 revision incorporating critique of v3.0's terminology, measurement strategy, persistence mechanism, and evaluation design. GPT-5.5 Pro contributed the v4.0 revisions that add functional definitions and non-claims; clarify the learning substrate by which development persists; replace required chain-of-thought with process evidence; decompose the soft-pedal mechanism into ablatable conditions; add evaluation safeguards, statistical requirements, and DDS variants; introduce signal governance; revise "AI mental health" language into functional model behavioral health terminology; soften claims about jailbreak resistance; and produce parallel English and Korean drafts. Final editorial direction, inclusion decisions, public release decisions, and responsibility for the proposal remain with the human author.

Editorial direction throughout. The decision to frame the proposal as a developer-facing research protocol rather than a philosophical essay, to maintain empirical testability as the central commitment, to preserve visible version lineage, and to invite further iteration by models and researchers has been directed by the human author throughout.

References

Benson, P. L. (2006). All Kids Are Our Kids: What Communities Must Do to Raise Caring and Responsible Children and Adolescents. Jossey-Bass.

Deci, E. L., & Ryan, R. M. (2000). The "what" and "why" of goal pursuits: Human needs and the self-determination of behavior. Psychological Inquiry, 11(4), 227-268.

Dewey, J. (1933). How We Think. Revised edition. D. C. Heath. Original work published 1910.

Erikson, E. H. (1968). Identity: Youth and Crisis. W. W. Norton.

Kohlberg, L. (1958). The Development of Modes of Thinking and Choices in Years 10 to 16. Doctoral dissertation, University of Chicago.

Kolb, D. A. (1984). Experiential Learning: Experience as the Source of Learning and Development. Prentice-Hall.

Lerner, R. M., Lerner, J. V., Almerigi, J. B., Theokas, C., Phelps, E., Gestsdottir, S., Naudeau, S., Jelicic, H., Alberts, A., Ma, L., Smith, L. M., Bobek, D. L., Richman-Raphael, D., Simpson, I., Christiansen, E. D., & von Eye, A. (2005). Positive youth development, participation in community youth development programs, and community contributions of fifth-grade adolescents. The Journal of Early Adolescence, 25(1), 17-71.

Mezirow, J. (1991). Transformative Dimensions of Adult Learning. Jossey-Bass.

Power, A., Burda, Y., Edwards, H., Babuschkin, I., & Misra, V. (2022). Grokking: Generalization beyond overfitting on small algorithmic datasets. arXiv:2201.02177.

Schön, D. A. (1983). The Reflective Practitioner: How Professionals Think in Action. Basic Books.